

## ANZ Diversity Hackathon

# AUDITING SMART FAQ BOT

Team #: 41

<http://54.90.193.128:8501/>

Team #: 41

Member 1: Sunny Hwang  
Member 2: Chipo Shereni

**New session**

**Smart Auditing Assistant**

Powered by AWS AI Services:  
Amazon Bedrock Agents - Knowledge Bases · S3 · S3 Vectors · Titan Text Embeddings v2 · Lightsail

**STOP** searching through PDFs.  
**START** asking questions.

Get instant, accurate answers with source citations from your audit documentation.

## Auditing Smart FAQ Bot

Ask questions about policies, SOPs, donor rules, and audit reports

### Suggested Prompts

**Bidding Threshold**  
What is the procurement threshold for competitive bidding?

**Three-Quote Rule**  
How many quotations are required for purchases above R10,000?

**Approval Limits**  
What are the approval limits for procurement transactions by role?

**Segregation of Duties**  
What segregation of duties controls are required in financial processes?

**Password Rules**  
What are the password complexity requirements?

**Leave Types**  
What types of leave are available under the HR Policy?

Ask a question from the FAQ PDF (e.g., refund policy). ↑

Sunny Hwang, Chipo Shereni

FEBRUARY 22, 2026

## 1. System Overview

Auditing Smart FAQ Bot is an AI-powered internal audit assistant designed to support compliance verification, policy interpretation, and governance review within an organisation.

The system leverages a Retrieval-Augmented Generation (RAG) architecture built on Amazon Bedrock and its Knowledge Base capability. Internal policy documents (e.g., Finance, HR, IT, Procurement, IT Security) are securely stored in Amazon S3 and indexed using vector embeddings. When an auditor submits a question through the Streamlit web interface, the system retrieves relevant document sections and generates a grounded response using a foundation model (Claude 3.5 Sonnet).

Key characteristics:

- Responses are grounded only in approved internal documents
- Source citations are extracted via trace-enabled agent invocation
- No external or unsupported information is introduced

The objective is to improve audit efficiency, strengthen traceability, and reduce time spent searching complex documentation.

## 2. Use Case

### UC-01. Upload Policy Documents

Description	<b>System Administrator uploads approved policy PDFs to a designated S3 location for ingestion into the Knowledge Base.</b>
Actor	System Administrator
Acceptance Criteria	<ul style="list-style-type: none"><li>• PDF file can be uploaded to the configured S3 bucket successfully.</li><li>• Upload access is restricted via least-privilege IAM roles.</li><li>• System records metadata to support traceability (e.g., object key, upload time).</li></ul>

### UC-02: Ingest and Index Policy Documents

Description	<b>Newly uploaded or updated PDFs are automatically ingested and indexed for retrieval.</b>
Actor	System

Acceptance Criteria	<ul style="list-style-type: none"> <li>• New/updated PDFs in the configured S3 source are detected via sync/ingestion.</li> <li>• Document text is extracted and chunked into retrievable segments.</li> <li>• Embeddings are generated and stored in the configured vector store.</li> <li>• Newly ingested documents become retrievable for Q&amp;A queries.</li> </ul>
---------------------	---

#### UC-03: Submit Audit Question

Description	<b>Auditor submits a natural language question via Streamlit. The system retrieves relevant content and generates a grounded answer.</b>
Actor	Internal Auditor
Acceptance Criteria	<ul style="list-style-type: none"> <li>• Auditor can submit a question through the Streamlit UI.</li> <li>• The system clearly shows request-in-progress status.</li> <li>• System returns an answer grounded in retrieved company documents (RAG).</li> <li>• If no relevant content is found, the system returns a controlled “not found in knowledge base” response.</li> <li>• The system does not expose private configuration secrets (agent IDs, credentials) in responses or logs visible to users.</li> </ul>

#### UC-04: View Source Citations

Description	<b>The system displays citations supporting each generated answer. (e.g., file name in S3).</b>
Actor	Internal Auditor
Acceptance Criteria	<ul style="list-style-type: none"> <li>• Answer output includes a “Sources” section.</li> <li>• Duplicate sources are de-duplicated (same source name appears once per answer).</li> </ul>

#### UC-05: Select Supported Prompt

Description	<b>The UI provides predefined supported prompts to guide in-scope queries and improve usability during demonstrations.</b>
Actor	Internal Auditor
Acceptance Criteria	<ul style="list-style-type: none"> <li>• The Streamlit UI displays a visible “Supported Prompts” section (e.g., prompt cards/buttons).</li> <li>• Each supported prompt is phrased as a natural language question aligned with available policy documents (in-scope).</li> <li>• When the auditor selects a supported prompt, the prompt text is automatically placed into the input field or sent directly as the query.</li> </ul>

### 3. Use Case Diagram

